Artificial Intelligence and Local Speech Synthesis

Chapter 1: The Evolution of Neural Voices
Neural text-to-speech models have achieved near-human expressive intonation. Kokoro-82M is an ultra-lightweight, 82-million parameter model capable of synthesizing studio-quality voices in real time on commodity hardware.

Chapter 2: Speech-to-Text Transcription
Faster-Whisper leverages CTranslate2 to transcribe audio up to four times faster than vanilla Whisper while utilizing significantly less memory.

Conclusion
This document demonstrates multi-page PDF text extraction and narration.